

# Python for Data Analysis Project

Clément MUFFAT-JOLY Jean MARCHAND – DIA4

---

## INTRODUCTION

---

We worked on the Dataset for estimation of **obesity levels** based on eating habits and physical condition in individuals from Colombia, Peru and Mexico.  
The information has 17 features and 2111 records.



# About the features of the dataset

The **attributes** related with **eating habits** are:

- Frequent consumption of high caloric food (FAVC)
- Frequency of consumption of vegetables (FCVC)
- Number of main meals (NCP)
- Consumption of food between meals (CAEC)
- Consumption of water daily (CH20)
- Consumption of alcohol (CALC)

The **attributes** related with the **physical condition** are:

- Calories consumption monitoring (SCC)
- Physical activity frequency (FAF)
- Time using technology devices (TUE)
- Transportation used (MTRANS)

The **targert** is NObesity with **different** values:

- Insufficient Weight
- Normal Weight
- Overweight Level I
- Overweight Level II
- Obesity Type I
- Obesity Type II
- Obesity Type III



# Import Dataset and EDA

---

We first **import** the dataset :

## 1. Import Dataset

```
df = pd.read_csv('ObesityDataSet.csv', encoding='latin-1')
```

Then we started to **explore** the dataset with the head function:

## 2 Exploration Dataset

```
# Preview of the dataset  
df.head()
```

We had to **convert** all **quantitative** features to int type:

```
df['Age'] = round(df['Age']).astype(int)  
df['FCVC'] = round(df['FCVC']).astype(int)  
df['CH2O'] = round(df['CH2O']).astype(int)  
df['FAF'] = round(df['FAF']).astype(int)  
df['TUE'] = round(df['TUE']).astype(int)
```

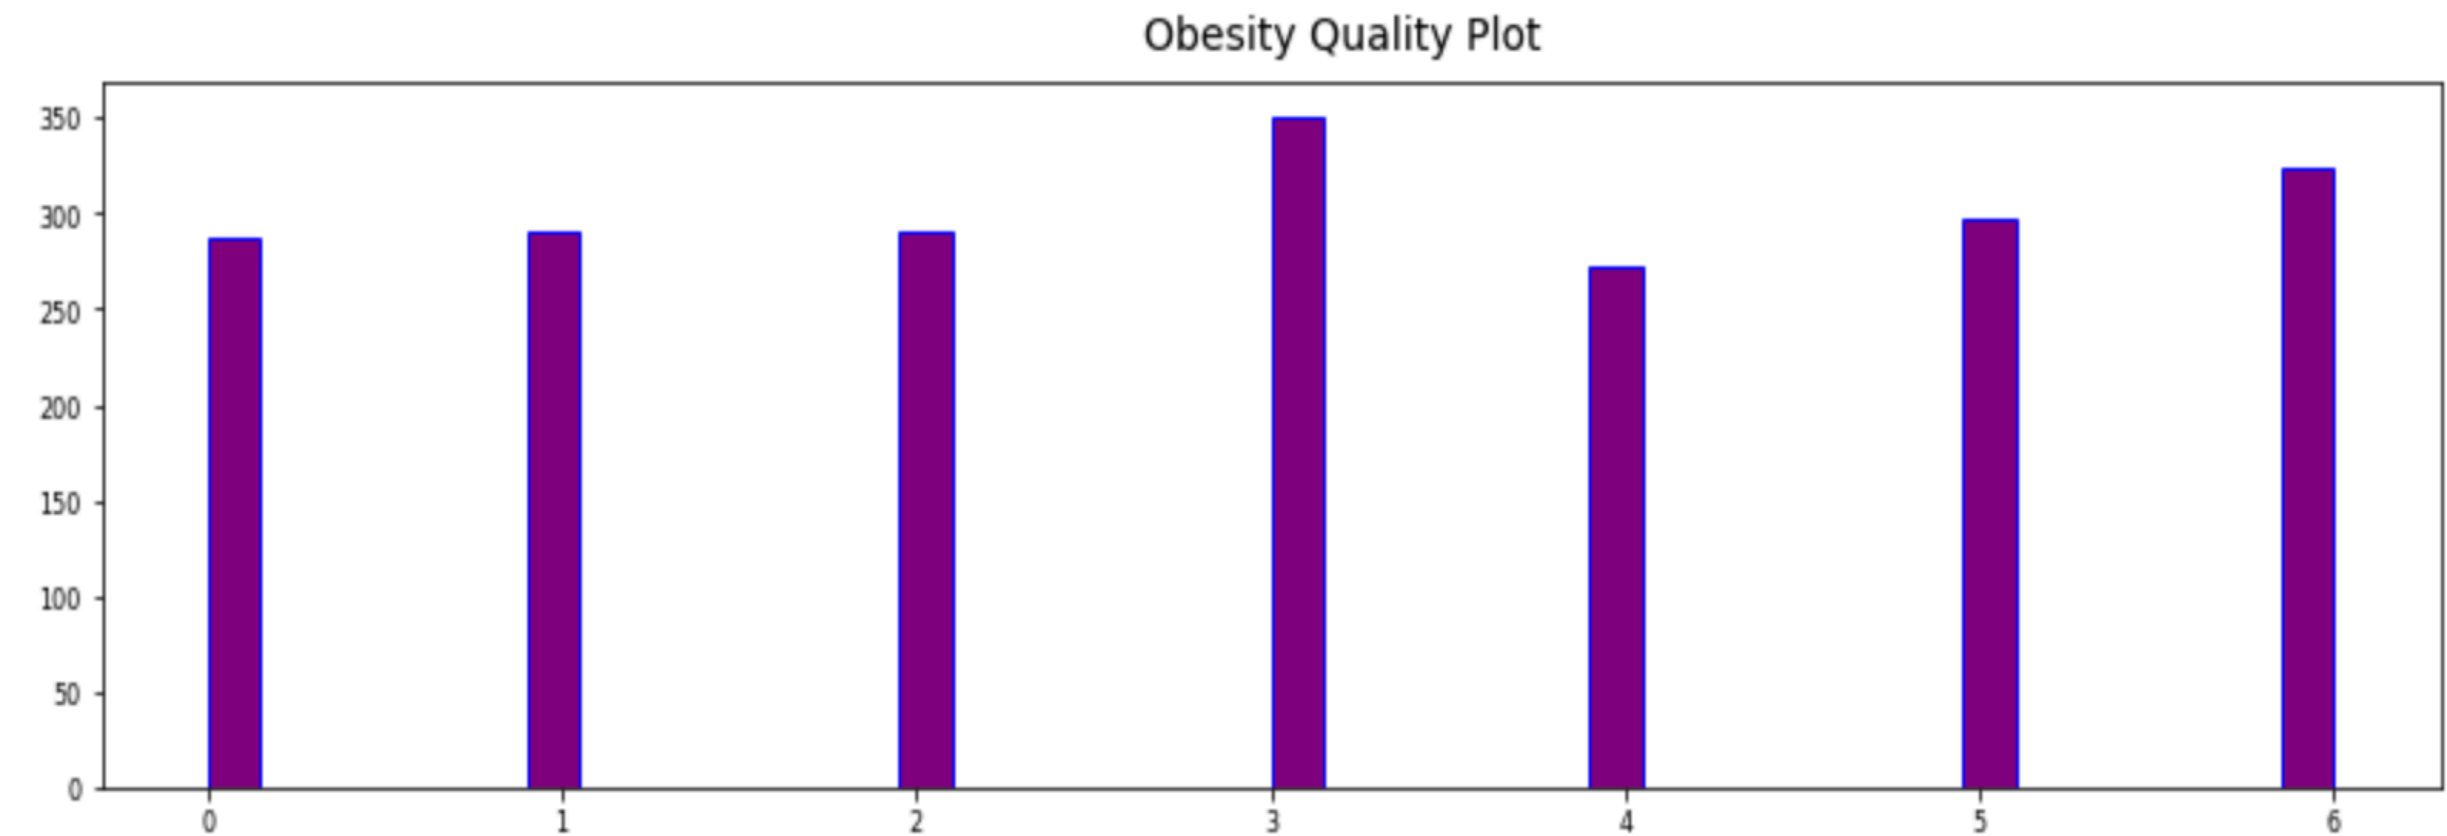
And we had to **convert** all **qualitative** features by replace with int values

## Replace qualitative features by numerical values


```
|: df["family_history_with_overweight"].replace(['yes', 'no'],[1,0],inplace=True)  
  
|: df["FAVC"].replace(['yes', 'no'],[1,0],inplace=True)
```

# Data Visualization


## Distribution of obesity categories




We have **explored** several data **visualizations** in the notebook. Here are 3 interesting graphs that we present to you. First, here is the **distribution** of the **number** of values (people) for **each level of obesity**. Since the qualitative values have been transformed into numbers, 0 corresponds to the most underweight people and 6 the most overweight. In 1 these are the so-called **normal weight** people. We can see that this category includes the **most people** is the 3 corresponds to overweight level 2. However, we can see that the more **extreme weight categories** are **very represented**, which is not a guarantee of good health for most of the people who participated in this survey. Let's try to study two other graphs to **better understand** these data




Bar chart



Stacked bar chart



Line graph




Gantt chart




Polar area diagram



Scatter plot



Calendar heatmap



Stacked area chart




Sparkline




Column sparkline

# Data Visualization


Distribution of obesity categories according to age




Bar chart



Stacked bar chart



Line graph




Gantt chart




Polar area diagram



Scatter plot



Calendar heatmap



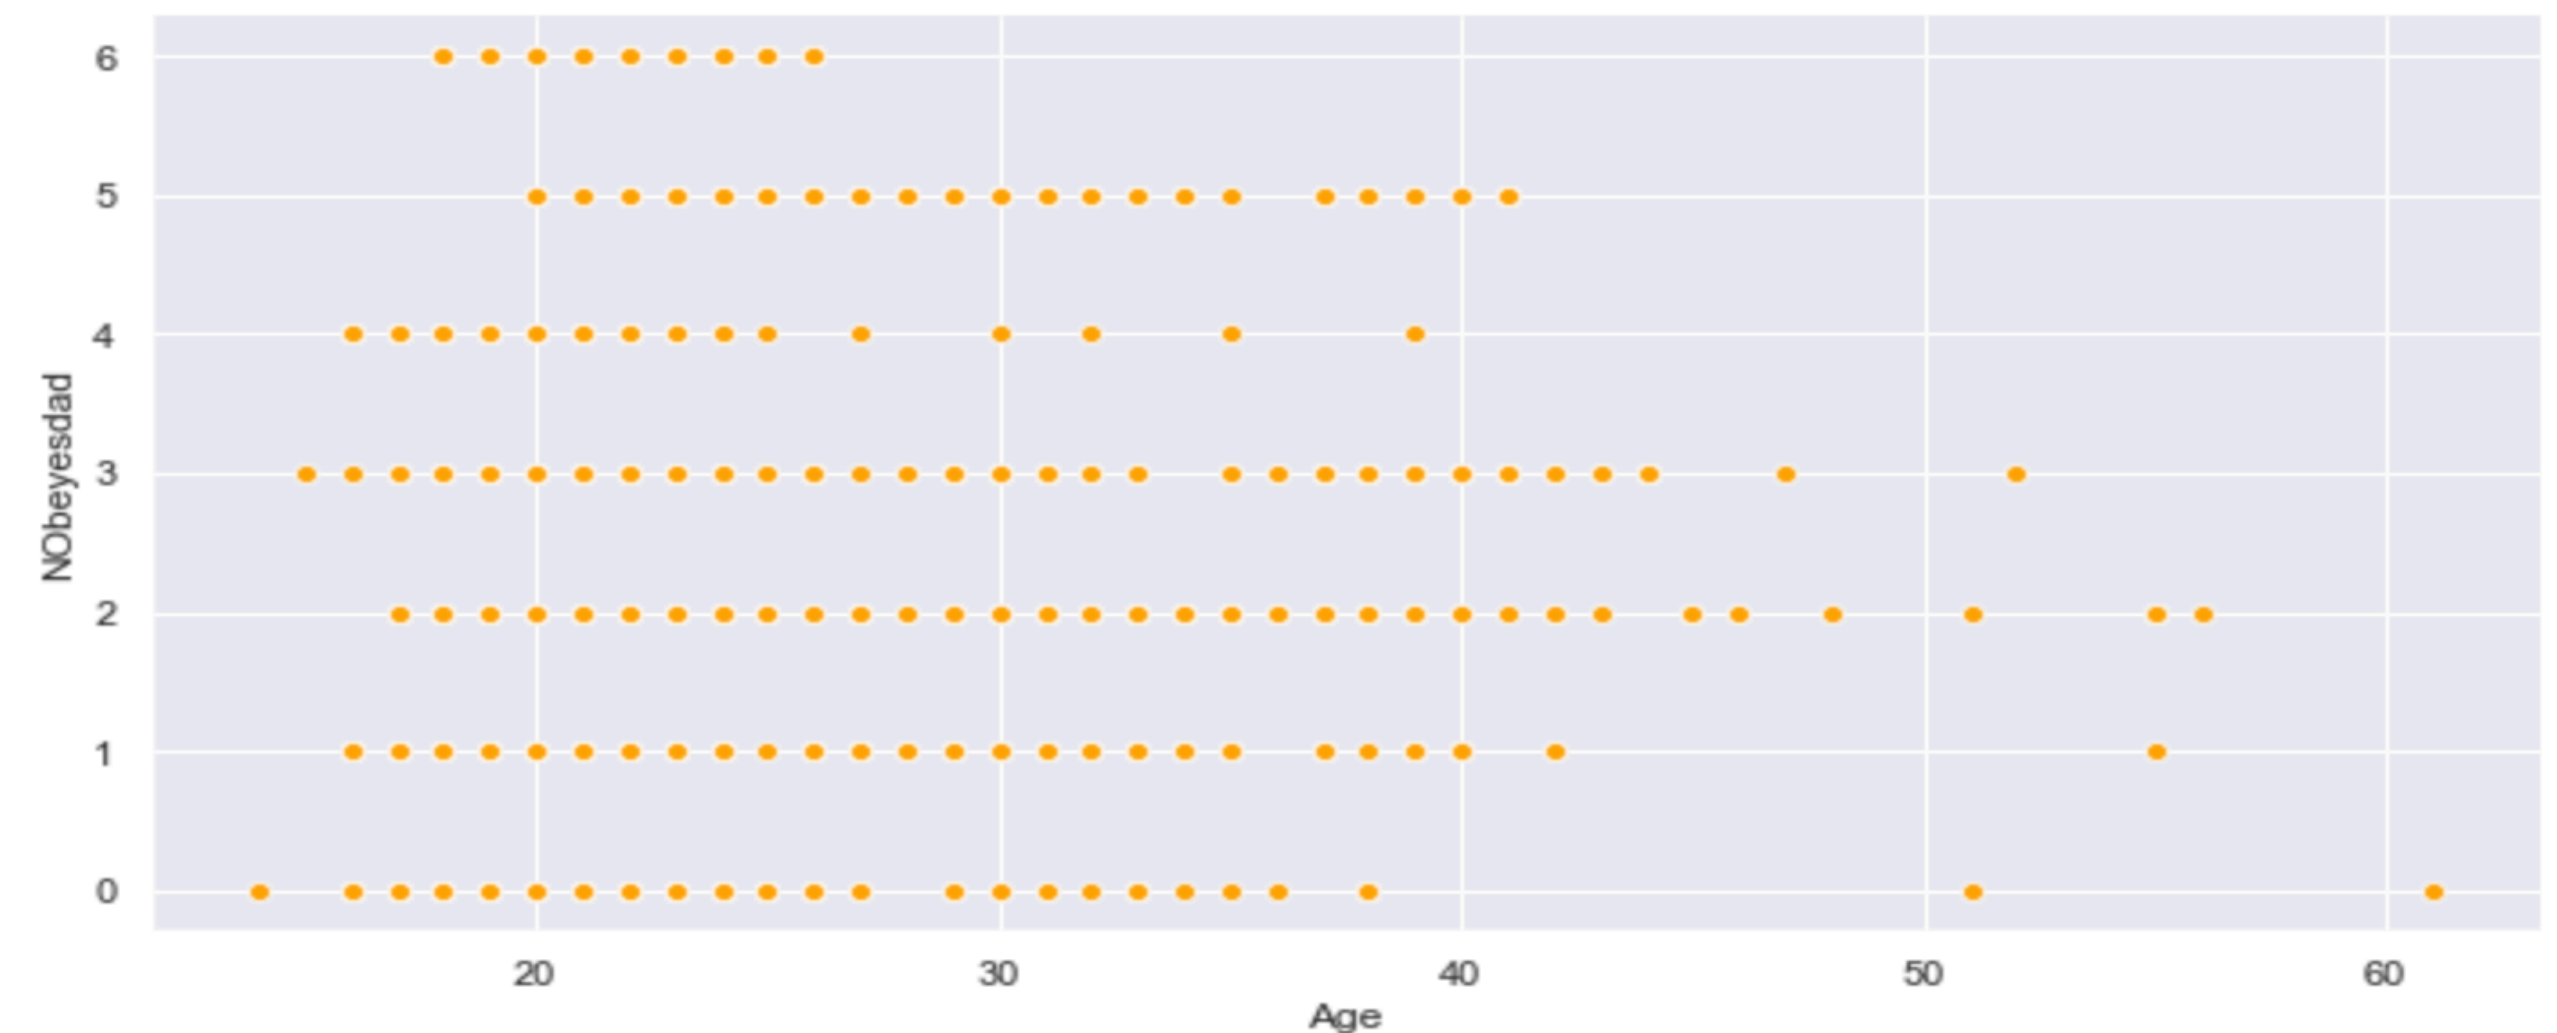
Stacked area chart



Sparkline



Column sparkline




In this second graph we can see how the **categories of obesity** are **distributed** according to **age**. At first glance the data are rather well distributed but we can observe that **no person over 26 years old** is at **level 3 of obesity** (the highest). We also note that **the majority of obesity levels 1 and 2 are made up of young people**. Conversely, the **oldest people** are divided between the **level of overweight 2** and the **level of underweight**. Is **age** therefore a **determining** variable in **the level of obesity**? This is what we will see next with the matrix of




# Data Visualization


## Percentage distribution of the population weight distribution




Bar chart



Stacked bar chart




Line graph




Gantt chart




Polar area diagram




Scatter plot



Calendar heatmap




Stacked area chart



Sparkline



Column sparkline



We can distinguish **3 bumps** on the curve which could imply that the population is **divided** into **3 large weight groups**. We can see that there are about **as many people** in the **underweight** part as **overweight**. An **interesting** data is the high distribution, about **15%** of the population towards an **average** weight of **75 kg** which is **more** than the **average weight** and confirms the **tendency** that this population is **overweight**.

# Pre-Processing

---

We **drop** features using the **calculation** of the predict value :

Then we **split** the data **without** 'Nobeyesdad' which is the **target** value.

We **scale** the value using **StandartScaler** function on **x\_train** and **x\_test**.

We **display** the **correlation map** with `sns.heatmap` to see the **correlation between features**

## Preprocessing data

Now, we can drop some values because they were used for the calculation of the predict value.

```
] : df.drop(['Height', 'Weight'], 1, inplace=True)
```

We split our data in training and test data:

```
] : x,y = df.loc[:,df.columns != 'Nobeyesdad'], df.loc[:, 'Nobeyesdad']  
    x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.3,random_state = 1)  
    print(df)
```

Now we can scale our data:

```
] : sc = StandardScaler()  
    x_train = sc.fit_transform(x_train)  
    x_test = sc.transform(x_test)
```


```
: #correlation map  
f,ax = plt.subplots(figsize=(18, 18))  
sns.heatmap(df.corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax).set(title="Correlation obesité")  
plt.show()
```

# Pre-Processing

## Correlation map

### Important features to predict obesity:

- We can see that the **family history of obesity** and **(FCVC) the frequency of vegetable consumption** are the **two** variables **most correlated** with the **level of obesity** (with a correlation of **0.3**). They are therefore **decisive**.
- **Next** comes a **correlation** of **0.2** (**FAVC**) for the **frequency of consuming high calorie foods**.
- And **finally** we have a correlation of **0.1**: **age**, the **(NCP) the number of meat-based meals**, **(CH20) the water consumption** per day and finally **(CALC) the alcohol consumption**.





# Machine Learning Models



---

In this part we try to **define** a **model** capable of **best predicting** the **level of obesity** based on certain variables.

We have tested **several models** that we will **present to you** one by one and you show at the end the **classification between them**.

We did :

- KNN
- Logistic Regression
- LDA
- Classification Tree
- Random Forest
- Bagging
- Boosting



# Machine Learning

## KNN


We fit the **KNN model** using **KNeighborsClassifier** from sklearn collection. We set to **3** the number of **neighbors**. We can see we have **0,73 accuracy**.

We chose to **display** the **accuracy** of **training\_set** and **test\_set** according to the **number of neighbors** defined in the knn. We can see that the **best model** is with **1 neighbor**, which gives an **accuracy of 0.75**.

### KNN

```
!]: # train test split
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(x_train,y_train)
prediction = knn.predict(x_test)
print('With KNN (K=3) accuracy is: ',knn.score(x_test,y_test)) # accuracy
```

With KNN (K=3) accuracy is: 0.7334384858044164



Best accuracy is 0.750788643533123 with K = 1

# Machine Learning

## Logistic Regression

We fit the **Logistic Regression model** using **LogisticRegression** from sklearn collection.


### Logistic Regression

```
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression(max_iter= 2000, random_state=10)
logreg.fit(x_train,y_train)
y_pred = logreg.predict(x_test)
logreg.score(x_train,y_train)
```

0.6161137440758294

Thanks to the **confusion matrix** we see that the **accuracy** of the model is **0.55**. It's **less** than the **KNN** model.

Accuracy Score : 0.556782334384858



# Machine Learning

## LDA

We fit the **LDA model** using **LinearDiscriminantAnalysis** from sklearn collection.

Thanks to the **confusion matrix** we see that the **accuracy** of the model is **0.34**. It's **less** than the **KNN** and **the Logistic Regression** model.

### LDA


```
: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

lda = LDA(n_components=1)
x_train_new = lda.fit_transform(x_train, y_train)
x_test_new = lda.transform(x_test)

classifier = RandomForestClassifier(max_depth=2, random_state=4)

classifier.fit(x_train_new, y_train)
y_pred = classifier.predict(x_test_new)
```

Accuracy0.34542586750788645



# Machine Learning

## Decision Tree Classifier

We fit the **Decision Tree model** using **DecisionTreeClassifier** from sklearn collection. We use the ‘*gini*’ criterion.

### Decision Tree Classifier


```
|: from sklearn import tree
   from sklearn.tree import DecisionTreeClassifier

   clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=0)

   # fit the model
   clf_gini.fit(x_train, y_train)
   y_pred_gini = clf_gini.predict(x_test)
```

Accuracy0.48580441640378547

Thanks to the **confusion matrix** we see that the **accuracy** of the model is **0,48**. It's **less** than the **KNN** and **the Logistic Regression** model but **better** than **LDA** model.





# Machine Learning

## Random Forest


We fit the **Random Forest model** using **RandomForestClassifier** from sklearn collection.

Thanks to the **confusion matrix** we see that the **accuracy** of the model is **0,81**. It's **better** than the **KNN**, the **Logistic Regression**, **LDA** and the **Decision Tree** model.

### Random Forest

```
: from sklearn.ensemble import RandomForestClassifier  
  
RF = RandomForestClassifier(n_estimators=300, random_state=50)  
RF.fit(x_train, y_train)  
y_pred = RF.predict(x_test)
```

Accuracy0.8091482649842271



# Machine Learning

## Bagging

We fit the **Bagging model** using **BaggingClassifier** from sklearn collection.

Thanks to the **confusion matrix** we see that the **accuracy** of the model is **0,81**. It's **better** than the **KNN**, the **Logistic Regression**, **LDA**, **Decision Tree** and the **Random Forest** model.


### Bagging

```
: from sklearn.ensemble import BaggingClassifier

bagging = BaggingClassifier(n_estimators=300, random_state=50)
bagging.fit(x_train, y_train)

y_pred = bagging.predict(x_test)
cm = confusion_matrix(y_test, y_pred)
```

Accuracy0.8123028391167192



# Machine Learning

## Boosting

---

We fit two **Boosting model**. The first one using **HistGradientBoostingClassifier** and the second using **GradientBoostingClassifier**.

Thanks to the **confusion matrix** we see that the **accuracy** of the two model is **0,80**. It's **better** than the **KNN**, the **Logistic Regression**, **LDA**, **Decision Tree**, **Random Forest** model but **less** than **Bagging**

### Histogram-based Gradient Boosting

```
] : from sklearn.experimental import enable_hist_gradient_boosting
    from sklearn.ensemble import HistGradientBoostingClassifier

    boost = HistGradientBoostingClassifier(random_state=50)
    boost.fit(x_train, y_train)
```

Accuracy0.804416403785489

### Gradient Boosting

```
] : from sklearn.ensemble import GradientBoostingClassifier

    gboost = GradientBoostingClassifier(n_estimators=300, random_state=50)
    gboost.fit(x_train, y_train)


    y_pred = gboost.predict(x_test)
```

Accuracy0.804416403785489

# Machine Learning

## Results

Percentage Of Success Of Algorithms And Test Result



We can see on this graph all the **accuracies scores** of **each** of the **models** on the **train\_set** and on the **test\_set**. We have decided to keep the **RandomForest** model which offers very **good results**.

# Machine Learning

## Features Selection

Before finalizing the model to be used on our API, we wanted to **simplify** it by **removing** the very **unimportant** variables in the role of predicting the level of obesity. So here are **the least important** values on the graph.

Finally we **keep** these variables:


- Gender
- Age
- Family\_history\_with\_overweight
- FCVC
- NCP
- CAEC
- CH2O
- FAF
- CALC

The **accuracy drops by 2%** but the model is simpler.

### Features selection

We have selected Random Forest algorithm because it is the best here.

```
|: indices = RF.feature_importances_.argsort()
plt.title('Feature Importances')
plt.barh(range(len(indices)), RF.feature_importances_[indices], color='b', align='center')
plt.yticks(range(len(indices)), [x.columns[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



We can keep only keep features with high importance (here we keep feature importance >0.05) to simplify our model:



# API

Framework used

---

To develop our **API** with python, we had **two choices: Django or Flask**. We chose **Flask** because it is **lightweight** than Django. And we didn't need complicated structures to our API, user just need to make a **single POST request** to obtain the prediction.





# API

## Website

---

Moreover, the **Flask** project is composed in **two parts**. One is the **API** that **any program** can **call**. And the second one is the **website** where **any user** can **interact** with to obtain a **prediction**. To predict the obesity level, we used the **model's dump** obtained by the **notebook**. Also, we used **Bulma** for the **CSS** of our website and **Pickle** to use the **model's dump** for the **website** and the **API**.

 GitHub



Obesity Analysis

What is your gender ?

☐ Male ☐ Female

What is your age ?

Do you usually eat vegetables in your meals ?

How many main meals do you have daily ?

Has a family member suffered or suffers from overweight ?

☐ Yes ☐ No

Do you eat any food between meals ?

How much water do you drink daily ?

How often do you have physical activity ?

How often do you drink alcohol ?

Make the prediction

Obesity Analysis by Jean Marchand and Clément Muffat-joly

 **BULMA**